

Ryan Shaw  
ryshaw

1.

```
def f3(s1, s2, y):

    assert y == 1 or y == 2
    # Forward pass: Compute loss
    L = None

    e1 = math.exp(s1)
    e2 = math.exp(s2)
    d = e1 + e2
    p1 = e1 / d
    p2 = e2 / d
    if y == 1:
        L = -math.log(p1)
    else:
        L = -math.log(p2)

    # Backward pass: Compute gradients
    grad_s1, grad_s2, grad_L = None, None, 1.0
    grad_p1, grad_p2 = None, None
    grad_e1, grad_e2 = None, None

    # Compute dL/dp1 and dL/dp2
    if y == 1:
        dp1 = -1 / p1
        dp2 = 0
    else:
        dp1 = 0
        dp2 = -1 / p2

    # Compute dp1/ds1, dp1/ds2, dp2/ds1, and dp2/ds2
    e1d = e1 / d
    e2d = e2 / d
    dp1ds1 = e1d * (1 - e1d)
    dp1ds2 = -e1d * e2d
    dp2ds1 = -e2d * e1d
    dp2ds2 = e2d * (1 - e2d)
```

```

# Compute dL/ds1 and dL/ds2 using the chain rule
grad_s1 = dp1 * dp1ds1 + dp2 * dp2ds1
grad_s2 = dp1 * dp1ds2 + dp2 * dp2ds2

grads = (grad_s1, grad_s2)
return L, grads

```

2.

(Fc-forward)

```

out = None
#####
# TODO: Implement the forward pass. Store the result in out. #
#####
out = np.dot(x, w) + b
#####
#                               END OF YOUR CODE           #
#####
cache = (x, w, b)
return out, cache

```

(fc-backward)

```

x, w, b = cache
grad_x, grad_w, grad_b = None, None, None
#####
# TODO: Implement the backward pass for the fully-connected layer #
#####
grad_x = np.dot(grad_out, w.T)
grad_w = np.dot(x.T, grad_out)
grad_b = np.sum(grad_out, axis=0)
#####
#                               END OF YOUR CODE           #
#####
return grad_x, grad_w, grad_b

```

(relu-forward)

```

out = None
#####
# TODO: Implement the ReLU forward pass. #
#####
out = np.maximum(0, x)
#####
#                               END OF YOUR CODE           #
#####
cache = x
return out, cache

```

(relu-backward)

```
grad_x, x = None, cache
#####
# TODO: Implement the ReLU backward pass. #
#####
grad_x = grad_out
grad_x[x <= 0] = 0
#####
#                                     END OF YOUR CODE #
#####
return grad_x
```

(softmax loss)

```
loss, grad_x = None, None
#####
# TODO: Implement softmax loss #
#####
# Shift x by max value to avoid numerical instability
x_max = np.max(x, axis=1, keepdims=True)
x_exp = np.exp(x - x_max)
x_sum = np.sum(x_exp, axis=1, keepdims=True)

# Compute softmax probabilities
p = x_exp / x_sum

# Compute loss
n = x.shape[0]
loss = -np.sum(np.log(p[np.arange(n), y])) / n

# Compute gradient of loss with respect to scores
grad_x = p.copy()
grad_x[np.arange(n), y] -= 1
grad_x /= n

#####
#                                     END OF YOUR CODE #
#####
return loss, grad_x
```

(l2-regularization)

```
loss, grad_w = None, None
#####
# TODO: Implement L2 regularization. #
#####
loss = 0.5 * reg * np.sum(w * w)
grad_w = reg * w
#####
#                                     END OF YOUR CODE #
#####
return loss, grad_w
```



(backward)

```
grads = {}
#####
# TODO: Implement the backward pass to compute gradients for all      #
# learnable parameters of the model, storing them in the grads dict  #
# above. The grads dict should give gradients for all parameters in  #
# the dict returned by model.parameters().                            #
#####
cache1, cache2, cache3 = cache
dX3, dW2, db2 = fc_backward(grad_scores, cache3)
dX2 = relu_backward(dX3, cache2)
dX1, dW1, db1 = fc_backward(dX2, cache1)
grads = {'W1': dW1, 'b1': db1, 'W2': dW2, 'b2': db2}
#####
#                               END OF YOUR CODE                       #
#####
return grads
```

4.



a. (training-step)

```
loss, grads = None, None
#####
# TODO: Compute the loss and gradient for one training iteration.      #
#####
#forward pass
scores, cache = model.forward(X_batch)

#calculate loss
loss, dloss = softmax_loss(scores, y_batch)
loss1, dloss1 = l2_regularization(model.parameters()['W1'], reg)
loss2, dloss2 = l2_regularization(model.parameters()['W2'], reg)
loss += loss1 + loss2

#backward pass
grads = model.backward(dloss, cache)
grads['W1'] += dloss1
grads['W2'] += dloss2
#####
#                               END OF YOUR CODE                       #
#####
return loss, grads
```

b. In your report, include the loss / accuracy plot for your best model, describe the hyperparameter settings you used, and give the final test-set performance of your model.





```
# How much data to use for training
num_train = 20000

# Model architecture hyperparameters.
hidden_dim = 300

# Optimization hyperparameters.
batch_size = 256
num_epochs = 100#500
learning_rate = 1e-2#6e-3
reg = 1e-6
```

- c. In your report, include the loss / accuracy plot for your overfit model and describe the hyperparameter settings you used.



```
# How much data to use for training
num_train = 2024

# Model architecture hyperparameters.
hidden_dim = 256

# Optimization hyperparameters.
batch_size = 256
num_epochs = 500#500
learning_rate = 5e-2#6e-3
reg = 0
```

5.

(architecture)

```
#downsampling
self.conv1 = nn.Conv2d(1,64,kernel_size=3, stride=2, padding=1) #14x14x64
self.pool = nn.AvgPool2d(kernel_size=2, stride=2) #7x7x64

#classifier
self.conv2 = nn.Conv2d(64,64,kernel_size=3, stride=1, padding=1) #7x7x64
self.conv3 = nn.Conv2d(64,16,kernel_size=3, stride=2, padding=1) #4x4x16 (stride=2 make 7->4, 64->16)
self.lin = nn.Linear(256,10)

#####
#                               END OF YOUR CODE                               #
#####

def forward(self, x):
    #####
    # TODO: Design your own network, implement forward pass here                #
    #####
    x = x.to(device)
    relu = nn.LeakyReLU()

    # Flatten each image in the batch
    x = self.conv1(x)
    x = relu(x)
    x = self.pool(x)

    x = self.conv2(x)
    x = relu(x)
    # # No need to define self.relu because it contains no parameters

    x = self.conv3(x)
    x = relu(x)

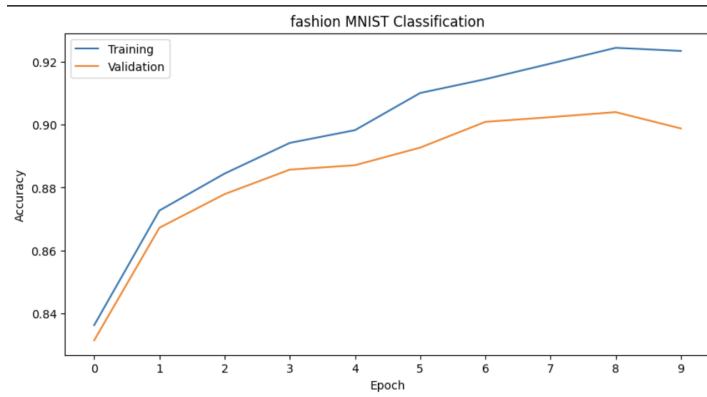
    x = x.view(x.size(0), -1)
    x = self.lin(x)
    # The loss layer will be applied outside Network class
    return x
```

(hyperparameters)

```
#####
batch_size=256
#####

# Set up optimization hyperparameters
learning_rate = 3e-3
weight_decay = 1e-5
num_epoch = 10 # TODO: Choose an appropriate number of epochs
optimizer = optim.Adam(model.parameters(), lr=learning_rate, weight_decay=weight_decay)
```

(plot)



(test set)

```
Evaluate on test set
100% |██████████| 40/40 [00:02<00:00, 16.35it/s]
Evaluation accuracy: 0.8961
CPU times: user 4min 56s, sys: 2.22 s, total: 4min 58s
Wall time: 5min 2s

0.8961
```

6.

(top image: good prediction; bottom: bad)

Predicted Class: Chihuahua



Predicted Class: sorrel

